

Peer Review: Local Multi-LLM Coding Workflow

Group A reviewing Group B's submission

Reviewers: Ahmad Alkaseb, Aleksandr Uktamovich Sorokin, Jeppe Jeppsson, Frederik Nicolajsen, Roman Ghazi, Kasper Amin

Strengths

- **Good explanation of the runs:** The report gives a clear account of what happened during the experiments.
- **Honest reporting:** The group acknowledges that documentation and deployment were not completed. This transparency is appreciated, even though those requirements remain unmet.
- **Failure analysis:** The discussion of hallucinated test success is valuable. It correctly highlights that an agent's claim that tests passed must be checked against actual test and build output.
- **Custom glue:** The report recognizes that pi.dev requires substantial custom code to recreate the agent team and orchestration.

Areas to improve

- **Different test projects:** The README-only test and the full 3D game differ greatly in complexity. The README test is useful for checking basic file operations, but the comparison should clearly distinguish that from end-to-end implementation capability.
- **Toolchain choice:** The report could better explain why OpenHands was selected and why the second toolchain, pi.dev, was rebuilt to reproduce the OpenHands agent team. OpenHands is more than a chatbot: it is a coding-agent environment with repository and command execution. Explain which capabilities justified the extra setup and which were easier to achieve in pi.dev.
- **Setup complexity:** Compare the practical setup and maintenance burden, including Docker/image patching for OpenHands versus the extensive custom extension, prompts, and configuration needed for pi.dev.
- **Endpoint verification:** The report clearly states which endpoint each role is configured to use, but does not verify which endpoint actually served each call. Configuration alone is not proof. Logging each call's model/endpoint—or checking loaded models with `ollama ps`—would strengthen the evidence.
- **Long runs and small models:** The 80-minute-plus runs did not produce a verified, working game. Smaller models are often better handled with short, incremental tasks, explicit context hand-offs, and Git checkpoints than with long runs that accumulate context and errors.
- **Failure-mode clarity:** Some explanations are hard to follow. For each failure, separate what happened, how it was detected, its impact, and how to prevent or recover from it.

Overall assessment

This is a thoughtful and candid evaluation, especially in its treatment of hallucinations and failed runs. The recommendation of pi.dev is plausible given its configurability, but the rationale could be stronger and the custom-code maintenance cost deserves emphasis. Neither toolchain delivered a verified runnable game, and endpoint routing was not independently confirmed. Clearer failure explanations and shorter, repeatable experiments would make the conclusions more convincing.