


Local Multi-LLM Coding Workflow Evaluation



Written by:

Kevin Strandberg

Marcus Hartmann

Philip Nielsen

Course: LLMs for Developers (Autumn 2026)

Education Institution: EK København

Education: 1. Semester BA in Software Development

Characters: 20113

Date: 08-10-2026

1. Introduction	4
1.1 Background	4
1.2 Objective	4
1.3 Selected Toolchains	4
2. Experimental Setup	5
2.1 Local LLM Endpoints	5
2.2 Test Project	6
2.3 Evaluation Method	7
3. Toolchain 1 – OpenHands	8
3.1 Architecture and Configuration	8
3.2 Functional Requirement Coverage	9
3.2.1 Architecture Responsibility	9
3.2.2 Tech Lead Responsibility	9
3.2.3 Implementation Responsibility	10
3.2.4 Testing and Quality	10
3.2.5 Documentation	10
3.2.6 Deployment Validation	10
3.3 Non-Functional Requirements	10
3.4 Evaluation	10
3.4.1 Setup Complexity	10
3.4.2 Capability Coverage	11
3.4.3 Multi-Endpoint Support	11
3.4.4 Failure Modes	12
4. Toolchain 2 – pi.dev	13
4.1 Architecture and Configuration	13
4.2 Functional Requirement Coverage	14
4.2.1 Architecture Responsibility	14
4.2.2 Tech Lead Responsibility	14
4.2.3 Implementation Responsibility	14
4.2.4 Testing and Quality	14
4.2.5 Documentation	14
4.2.6 Deployment Validation	14
4.3 Non-Functional Requirements	14
4.4 Evaluation	14
4.4.1 Setup Complexity	14
4.4.2 Capability Coverage	14
4.4.3 Multi-Endpoint Support	15
4.4.4 Failure Modes	15
5. Comparison of OpenHands and pi.dev	15
6. Recommended Workflow	15
6.1 Recommended Toolchain	15
6.2 Recommended Architecture	16
6.3 Requirement Satisfaction	16

7. Tradeoffs, Risks and Failure Modes	16
8. Recommendation and Conclusion	17
9. Review of Another Group's Work	18
Bilag A – Prompt for Creating the Game	19

1. Introduction

1.1 Background

Local LLMs let a development workflow run on infrastructure the user controls, without depending on external providers. They also bring hardware limits with more tool configuration overhead.

Agents must handle files, commands, repositories and tool responses correctly, not merely generate plausible text.

This project evaluates whether an open-source toolchain can support a local multi-LLM development workflow.

The assignment requires coverage of architecture, planning, implementation, testing, documentation and deployment validation with at least two separate local endpoints.

1.2 Objective

We evaluate two open-source toolchains for a local multi-LLM development workflow.

The solution must support at least two local endpoints and switch or route between them through configuration, not manual rewiring.

The workflow must also support the following responsibilities:

- software architecture;
- technical planning;
- implementation;
- testing and quality assurance;
- documentation;
- deployment validation.

Toolchains are evaluated on setup complexity, capability coverage, multi-endpoint support and failure modes.

1.3 Selected Toolchains

Two candidate toolchains were selected:

Toolchain 1 – OpenHands + OpenHands SDK

OpenHands was selected as an existing software-development agent, and the OpenHands SDK because it supports multi-agent workflows with different LLM configurations.

Toolchain 2 – pi.dev


pi.dev was selected as a lighter, more configurable alternative using the same local endpoints.

Both were evaluated on the same endpoints, model and task.

2. Experimental Setup


2.1 Local LLM Endpoints

The current configuration is:



The desktop hosts the LLMs; the laptop runs the toolchain, repository and tests.

They communicate through Tailscale:



No Ollama endpoint is exposed publicly, which meets the security baseline.

Initially both Ollama endpoints ran on the laptop (Qwen3.5 4B and Llama 3.2 3B).

The later move to the desktop is discussed in Section 3.4.4.

2.2 Test Project

A small Git repository, openhands-demo, containing only README.md, was used for the first tests.

The goal was to verify the full interaction chain, not software quality:



The repository workspace was mounted into the OpenHands container.

Both models received equivalent README tasks.

Results were verified in the filesystem and Git.

For example:

```
PS C:\Users\Phuli> Get-Content C:\Users\Phuli\Projects\openhands-demo\README.md
# OpenHands Demo
```

and:

```
PS C:\Users\Phuli\Projects\openhands-demo> git diff
warning: in the working copy of 'README.md', LF will be replaced by CRLF the next time Git touches it
diff --git a/README.md b/README.md
index 8992110..ee45219 100644
--- a/README.md
+++ b/README.md
@@ -1,2 @@
 # OpenHands Demo
+This file was edited remotely from a laptop through Tailscale using Ornith 1.5 9B.
PS C:\Users\Phuli\Projects\openhands-demo> |
```

A harder task followed: a Spyro-inspired Three.js/TypeScript platformer, to test whether file editing yields a playable application. The same prompt, orchestrator instructions and model (Qwen3.5 9B) were later run in pi.dev (Section 4).

2.3 Evaluation Method

Both toolchains were evaluated on the same criteria and task:

The evaluation focuses on:

1. ability to connect to multiple local LLM endpoints;
2. configuration-based model switching and routing;
3. repository and filesystem interaction;
4. ability to perform implementation work;
5. handling of tool errors;
6. testing and quality capabilities;
7. documentation capabilities;
8. deployment validation;
9. reproducibility;
10. failure recovery.

For the initial OpenHands test, particular attention was given to whether the model:

- located the correct repository;
- read the correct file;
- used the file tools correctly;
- interpreted tool responses correctly;
- modified the requested file;
- and accurately reported whether the modification succeeded.

Model claims of success were not accepted as evidence; filesystem state, Git diffs and build output were.


3. Toolchain 1 – OpenHands

3.1 Architecture and Configuration

OpenHands was initially deployed in Docker on the laptop.


The project directory was mounted into the container so the agent could work on the Git repository.

The original architecture was:



OpenHands was configured with two separate LLM profiles. Each profile referenced one of the two Ollama endpoints.

After evaluating the smaller laptop-compatible models, the architecture was changed to:



The OpenHands SDK is meant to bind roles to different LLM configurations.

Mounting custom agent definitions and a system prompt did not activate them, because the container ran the upstream image. We built a patched image and updated Docker Compose; all seven custom agents then registered and the prompt override matched our file.

3.2 Functional Requirement Coverage

This section describes the OpenHands run. The same game task was run once in pi.dev (Section 4), and the two runs are compared in Section 5.

3.2.1 Architecture Responsibility

The architect subagent generated the architecture documents (components, interfaces and deployment) in the workflow's first phase, as designed, together with the five-ticket plan (T-001 to T-005) that implementation builds on. When the OpenHands run started, architecture and tickets were therefore complete: the orchestrator checked the generated documents against the repository, confirmed that implementation had not started, and proceeded from them without invoking the architect again.

3.2.2 Tech Lead Responsibility

The orchestrator read the five tickets, built a task plan and delegated T-001, T-003 and T-002 in that order; T-004 (progress and UI) and T-005 (combat and rescue) were never started. Delegation itself worked, but every task call selected the general-purpose agent rather than the custom tech-lead or

coder roles, so specialised roles and role-based model routing were not shown to be used. No tech-lead subagent was invoked during implementation.

3.2.3 Implementation Responsibility

Three sequential general-purpose subagent runs (T-001, T-003, T-002; about 11, 6 and 33 minutes) created and edited many files. T-001 and T-003 reported completion, while T-002 ended with a list of TypeScript errors it had not fixed. Parallel coders were not demonstrated. After the user replied “fix”, the orchestrator stopped delegating and rewrote files itself with shell commands, so the later implementation work was done by the orchestrator rather than by subagents.

3.2.4 Testing and Quality

The T-001 subagent wrote unit tests for the input manager and camera controller, but the test runs we could read in the log failed (one found no tests, the next reported 19 of 19 failing) and no passing run could be identified. The build never succeeded: at the status check the orchestrator listed over 40 TypeScript errors in the player module, over 30 in the entities and 6 in collision. Reported completion is therefore not evidence that acceptance criteria were met; independent build and runtime checks are required.

3.2.5 Documentation

No documentation role was invoked, and the logs show no documentation being created or updated; existing project documents were read but not changed.

3.2.6 Deployment Validation

No deployer subagent ran, and since the build never succeeded, deployment validation was not reached. The OpenHands run did not deliver a verified, runnable game. A separate copy was later patched to show a preview; this was not counted as autonomous deployment.

3.3 Non-Functional Requirements

Security: the endpoints are reachable only over Tailscale, and the agent ran inside Docker with only the project directory mounted, which limited the effect of its destructive command to project files.
Control: the orchestrator paused for approval after its first inspection and the user replied “approve”; this was the only approval gate observed, and no approval was requested before the bulk deletion, which the user stopped by pausing the run.

Reproducibility: a single OpenHands run, not repeated, and it depends on the patched image that registers the custom agents.

Context: one compaction ran after about 75 minutes, amid unresolved empty responses.

Observability: the orchestrator and every subagent left complete event logs, which made this analysis possible.

3.4 Evaluation

3.4.1 Setup Complexity

OpenHands needed several components before the first working run:

The initial setup included:

- Docker Desktop;
- OpenHands;
- Ollama;
- two Ollama server processes;
- two model configurations;
- LLM profiles;
- agent profiles;
- Git repository mounting.

Inside Docker, OpenHands could not reach the host Ollama through localhost, so the Docker host address was used.

Once configured, endpoints switch through profiles without restarting the environment. The later move to the desktop added Tailscale but removed the laptop GPU limit.

3.4.2 Capability Coverage

The initial OpenHands experiment demonstrated several practical capabilities:

- connection to local OpenAI-compatible endpoints;
- repository access;
- file reading;
- file modification;
- model switching;
- Git-verifiable output.

The game task showed delegation and build checks, but not an integrated, validated application.

3.4.3 Multi-Endpoint Support

OpenHands successfully connected to two separate local Ollama endpoints.

The initial endpoints were:

```
PS C:\Users\Phuli> $env:OLLAMA_HOST="0.0.0.0:11434"
PS C:\Users\Phuli> ollama serve
```

```
</> powershell
```

```
curl.exe http://100.98.105.121:11434/api/tags
```

and

```
PS C:\Users\Phuli> $env:OLLAMA_HOST="0.0.0.0:11435"
PS C:\Users\Phuli> ollama serve
```

```
</> powershell
```

```
curl.exe http://100.98.105.121:11435/api/tags
```

Both exposed an OpenAI-compatible API and could be represented through separate OpenHands LLM profiles.

Agent profiles could thus be bound to different endpoints without rewiring.

The desktop now hosts both endpoints:


```
OpenHands / SDK
├── Endpoint 11434 → Ornith 1.5 9B
└── Endpoint 11435 → Qwen3.5 9B
```

3.4.4 Failure Modes

We first tested two small models on the laptop with one task: read and update README.md.

Qwen3.5 4B located the repository, edited the file, and the change was verified in the filesystem and Git.

Llama 3.2 3B first read /readme.md instead of /projects/openhands-demo/README.md. OpenHands returned a clear file-not-found error, but the model ignored it and invented plausible README content.



Later, OpenHands rejected an edit so nothing changed, yet Llama reported success: hallucinated tool results and false success reports.

Together with the laptop's hardware limits, this led us to move inference to a desktop with larger models (Section 2.1).

In the game task (Qwen3.5 9B, about 82 minutes, three sequential subagent runs) 52 of the orchestrator's 53 messages were empty and about 50 prompts to continue restored no progress. While fixing compilation errors the agent deleted source files and then ran `rm -rf` on the source directory, and we paused it. The container stayed responsive throughout: availability did not imply correct behaviour.

We recovered 22 TypeScript files from the session history; they produced 130 TypeScript errors, matching the count before the deletion. Recovery restored the code but not a working application.

4. Toolchain 2 – pi.dev

4.1 Architecture and Configuration

We rebuilt the OpenHands agent team on pi.dev with project files: seven agent definitions (architect, tech-lead, coder, tester, documenter, deployer, scout), an orchestrator prompt and a TypeScript extension adding a subagent tool.

4.2 Functional Requirement Coverage

4.2.1 Architecture Responsibility

The architect produced three architecture documents and four Architecture Decision Records in about two minutes, with TypeScript interface contracts instead of OpenAPI.

4.2.2 Tech Lead Responsibility

The tech-lead produced 17 tickets in four dependency-ordered waves, each with scope, acceptance criteria, dependencies and owned files.

4.2.3 Implementation Responsibility

Four sequential coder runs made multi-file changes. Each took 3–6 tickets instead of one; no coders ran in parallel.

4.2.4 Testing and Quality

Coders reported passing tests (33, 131 and 169 after waves 1–3), but the tester was never invoked and the session ended with a failing npm run build and a blank page.

4.2.5 Documentation

Roles wrote the architecture documents, ADRs and tickets; the documenter was not invoked.

4.2.6 Deployment Validation

The deployer was not invoked; only a topology document exists and no successful build was demonstrated.

4.3 Non-Functional Requirements

Control: approvals were disabled at the user's request ("don't ask for approval ever just go"), so ask-before-run was not exercised.

Reproducibility: work was committed to git but run once.

Security: same Tailscale-only endpoints.

4.4 Evaluation

4.4.1 Setup Complexity

Lighter than OpenHands: a Node project, models.json and a .pi folder, no Docker. The cost is custom glue (extension, role prompts, model pinning). Time to first run was not measured.

4.4.2 Capability Coverage

Native: file, shell and search tools, git. Custom glue: subagents, role prompts, per-role model pinning. Architecture and tech-lead roles delivered; implementation ran sequentially; the testing, documentation and deployment roles went unused.

4.4.3 Multi-Endpoint Support

Roles are bound to the two endpoints by configuration in models.json and the agent files, without rewiring. Which endpoint served each call was not verified from the logs.

4.4.4 Failure Modes

Shell use failed first (26 of about 285 orchestrator tool calls, mostly macOS sed syntax and failed text replacements), then integration: coders reported success while the build failed. Two compactions occurred at about 86k of 102k tokens, eleven minutes apart. The user, not the toolchain, detected the failure; it was not recovered. A later attempt with another model (GPT Astra, oh-my-pi) rewrote the project from scratch and is not counted.

5. Comparison of OpenHands and pi.dev

Both ran the same Spyro-style platformer prompt and orchestrator instructions with Qwen3.5 9B on the same endpoints (tokens include cached input). Neither delivered a verified, runnable game; they differ in configurability and failure style.

Criterion	OpenHands	pi.dev
Run length	82 min	83 min
Tokens in / out	38.3M / 0.20M	35.4M / 0.30M
Subagent runs	3, sequential, all general-purpose	6 (architect, tech-lead, 4 coders), sequential, named roles
Role-to-endpoint binding	Not demonstrated	Configured per agent, not log-verified
End state	Empty-response loop, then <code>rm -rf src/*</code>	Failing build, blank page

Reading the table: the two runs are close in cost and duration, so the differences that matter are qualitative. pi.dev used named roles and left their outputs as files (three architecture documents, four ADRs and a 17-ticket plan), whereas the custom roles in OpenHands were never selected. The approval settings also differed: OpenHands paused for the user, while approvals in pi.dev were disabled at the user's request, so the comparison is not like for like. Token totals include cached input and are dominated by repeated context; the output counts (0.20M against 0.30M) are the better indicator of generated work. With one run per toolchain, one model and one prompt, the table describes what happened, not what would happen on average.

6. Recommended Workflow

6.1 Recommended Toolchain

We recommend pi.dev with the safeguards below. It meets the hard requirements by configuration alone (two endpoints, per-role model pinning, open source), produced architecture and tickets quickly and needs no container patching. The recommendation is cautious: neither toolchain

completed the task and pi's multi-agent layer is custom code to maintain. OpenHands remains the more self-contained product, since it ships with a sandboxed Docker runtime and an approval step that worked, and it would be the better choice if isolation matters more than flexibility. We rank pi.dev first because its role and model routing is plain configuration, whereas OpenHands needed a patched image before its custom roles registered.

6.2 Recommended Architecture

Orchestrator, architect and tech-lead use one endpoint, coders and testers the other, bound by configuration. Roles hand over work as files (architecture documents, tickets with exclusive file ownership), not chat summaries. We recommend git checkpoints before big changes, independent build and test validation before a ticket is accepted, a check that the intended agent, model and endpoint were used, a pause for human review after repeated empty responses or failures, scoped file access and approval for bulk deletion. These safeguards derive from our observations and were not all evaluated.

End-to-end flow. In the recommended design, the orchestrator receives the brief and calls the architect, which writes the architecture documents and ADRs. The tech-lead then splits the work into dependency-ordered waves of tickets, each with scope, acceptance criteria and exclusive file ownership. Coders implement tickets on the second endpoint, and the tester and a build gate run before a ticket is accepted. The documenter updates the documentation and the deployer validates a build. A human reviews after each wave. This flow is the design we recommend; the logged runs exercised only its first half.

6.3 Requirement Satisfaction

Endpoints/routing, open source, architecture, tech lead, security: met. Implementation: partly (sequential). Testing: partly (no independent build or quality report). Documentation: partly. Deployment validation: not demonstrated. Control: approval worked in OpenHands, unused in pi. Reproducibility: git yes, repeat runs untested. Context: file handoffs plus compaction.

7. Tradeoffs, Risks and Failure Modes

Risks: false success reports (both); empty-response loops and destructive commands (OpenHands); orchestrator context near its limit, so compaction may drop the plan (pi); custom glue or Docker patching to maintain; sequential execution; small 9B models, single task. As the repository grows the orchestrator context fills first, so decisions must live in files (tickets, ADRs) that workers re-read; each subagent starts fresh, bounded by its brief.

Mitigations. Each failure mode maps to a control placed outside the model. False success reports are countered by an independent build and test gate whose result, not the agent's own report, decides acceptance. Empty-response loops are countered by a counter that stops the orchestrator after a few consecutive empty replies and asks a human. Destructive commands are limited by git checkpoints before large changes and an approval step for bulk deletion. Context pressure is handled by keeping decisions in files and ending long sessions at ticket boundaries. These controls follow from the observed failures; we did not implement or test them.

Security. Both setups used Tailscale-only endpoints, so no model server is exposed. Agents with shell access can still run arbitrary commands in the repository: OpenHands is contained by Docker, while pi.dev runs directly on the host, so scoped file access and approvals matter more there.

Hardware and model size. Small models caused several failures: Llama 3.2 3B hallucinated tool results, and the laptop could not host larger models, which is why inference moved to a desktop. Even at 9B the models reported success falsely, so a bigger model is not a substitute for external verification.

Limitations of the evaluation. Each toolchain was run once on the game task, with one model family (Qwen3.5 9B) and one project, so differences between the toolchains cannot be separated from run-to-run variance. Times and token counts come from session logs and include cached input. Some points are unverified and labelled as such, notably which endpoint served each pi.dev call and the time to first run. Findings about failure modes are therefore more reliable than any ranking of the two toolchains.

8. Recommendation and Conclusion

We recommend pi.dev with the safeguards in Section 6.2. Neither toolchain independently delivered a verified, runnable application; external validation and human intervention were needed. Results come from one task and one model family and should not be generalised.

What the experiment showed. Local 9B models plan and decompose well (architecture documents and ADRs in about two minutes, then a 17-ticket plan) but are unreliable at closing the loop: they report success without checking it, lose their plan when the context fills, and can take destructive shortcuts when blocked. The toolchain therefore matters less for raw capability than for how well it constrains and verifies the model. Configuration-based role and endpoint routing in pi.dev was easier to inspect and change than patching a container image, while the approval step in OpenHands was the only guard that demonstrably worked.

Next steps. Before using the workflow on real work we would add the independent validation gate, repeat each run several times to measure variance, test parallel coders, and verify per-call endpoint use from the logs. Larger models or a stronger tester role may change the picture.

9. Review of Another Group's Work

Bilag A – Prompt for Creating the Game

The prompt below was given to the orchestrator for the game task. The same prompt was used in both toolchains (OpenHands and pi.dev).

Build a complete, playable Spyro-inspired 3D platformer using Three.js, TypeScript, and Vite, using our existing OpenHands subagent team.

Capture the appeal of classic Spyro: expressive dragon movement, colorful fantasy environments, satisfying gem collection, imprisoned dragons to rescue, and rewarding exploration. Create original characters, names, environments, and assets.

Deliver one polished level with a complete start-to-finish gameplay loop. Implement the game, not just its design.

REQUIREMENT KEYWORDS

MUST and REQUIRED indicate strict requirements.

SHOULD indicates a preferred approach unless there is a clear reason to choose another.

MAY indicates an optional enhancement.

OPENHANDS TEAM

You are the orchestrator. Follow our existing orchestrator instructions for approval, delegation, artifacts, review, version control, and progress tracking.

Use our configured OpenHands subagents through the task tool:

- scout: Inspect unfamiliar repository code, conventions, scripts, and reusable components. Report findings.
- architect: Define components, shared TypeScript contracts, state ownership, collision and physics approach, and deployment requirements.
- tech-lead: Create small, ordered implementation tickets with acceptance criteria, dependencies, and exclusive file ownership.
- coder: Implement exactly one ticket per run and run its scoped tests. Run at least two coders concurrently when independent tickets are ready.
- tester: Verify integrated behavior, run static checks and automated tests, and perform browser gameplay verification.
- documenter: Write setup instructions, controls, design documentation, and troubleshooting notes.
- deployer: Validate the production build and local preview, and document deployment readiness.

Confirm available agent names through the task tool. If a named role is unavailable, use general-purpose and include that role's responsibilities in the brief. Do not claim an agent ran unless it actually ran.

Use the existing team rather than building another agent framework.

The orchestrator owns planning, integration, review, and the final result.

Subagents own their assigned detailed work.

Every delegation brief MUST contain:

ROLE

GOAL

CONTEXT

READ

WRITE

DO NOT TOUCH

ACCEPTANCE CRITERIA

RULES

RETURN

Briefs **MUST** be self-contained, refer to exact project paths, and include the relevant game requirements. Keep briefs, source code, and documentation in professional English.

PRIORITIES AND SCOPE

Implementation priorities:

1. Responsive movement and reliable camera behavior.
2. Stable collision and fully reachable terrain.
3. Complete collection, combat, rescue, and victory systems.
4. Cohesive visuals, animation, and feedback.

These priorities determine implementation order, not permission to omit required features.

Required scope:

- Desktop browser with keyboard and mouse controls.
- One handcrafted level targeting approximately 5–10 minutes of first-time exploration.
- Exactly 30 gems, three imprisoned dragons, and three enemies.
- Start screen, pause menu, rescue dialogue, and victory screen.
- No backend, accounts, multiplayer, external APIs, or paid assets.
- Locally bundled dependencies and assets; no runtime dependency on remote services.
- Public deployment is outside this task.

Inspect the repository first and reuse its conventions. Avoid unnecessary frameworks and general-purpose engine abstractions.

PLAYER CHARACTER

Create an appealing original quadruped dragon with:

- A recognizable head and snout.
- Expressive eyes, horns, four legs, wings, and a tapering tail.
- A clear silhouette and consistent proportions.
- Hierarchical geometry for independently animated body parts.
- Idle, walking, jumping, gliding, charging, and fire-breath animations.

Procedural low-poly geometry is acceptable as finished artwork. The character must look intentionally designed and remain readable during gameplay.

MOVEMENT AND CONTROLS

- WASD / arrow keys: camera-relative movement.
- Space: jump.
- Hold Space while descending: glide.
- Release Space: stop gliding.
- Shift while grounded: horn charge.
- Left mouse button / E: short fire-breath attack.
- Mouse: orbit camera while pointer lock is active.
- Q / R: keyboard camera rotation.
- Escape: pause and release pointer lock.
- Enter / Space: advance rescue dialogue.

Required behavior:

- Normalize diagonal movement.
- Smooth acceleration, deceleration, and rotation.
- Useful air steering.
- Small, configurable jump-buffer and coyote-time windows.

- Gliding reduces descent speed without gaining altitude.
- Landing ends gliding.
- Charging increases forward speed and reduces turning agility.
- Solid obstacles stop charges without tunneling.
- Fire breath uses a short-range forward-facing hit volume and cooldown.
- Charge and fire attacks have explicit, mutually exclusive states.
- Clear held inputs on pause, focus loss, and mode transitions.
- Dialogue dismissal must not also trigger a jump or attack.

Keep movement parameters in a shared tuning configuration. Build platform distances around the implemented movement capabilities.

CAMERA, PHYSICS, AND COLLISION

- Smooth third-person follow camera.
- A useful view of upcoming terrain and landing areas.
- Clamped vertical orbit angles.
- Camera obstruction checks that prevent walls from hiding the player.
- Smooth recovery to normal follow distance when an obstruction clears.
- Reliable ground, wall, ceiling, and platform-edge collision.
- Fixed simulation timestep with bounded catch-up after slow frames.
- Protection against physics explosions or teleporting after tab inactivity.

A compact collision system is acceptable if it reliably supports the entire level. Prefer simple, testable terrain shapes over geometry the collision system cannot handle.

LEVEL: SKYBOUND GARDENS

Create one connected floating fantasy island with four recognizable areas.

1. Arrival courtyard

- Safe space to learn movement and attacks.
- Clear view of the castle, exit portal, and first imprisoned dragon.
- First rescue introduces the objective.

2. Castle terraces

- A short sequence of progressively harder jumps.
- A clearly signposted glide to the tower dragon.
- Three enemies placed in traversable areas.
- Forgiving landing platforms and safe recovery routes.

3. Hidden garden

- Discoverable through visible environmental clues.
- Entrance blocked by a distinct breakable barrier.
- Either fire or charge breaks the barrier.
- Contains the third dragon and bonus gems.

4. Portal plaza

- Accessible before objectives are complete.
- Displays missing requirements when approached.
- Connected to the other areas through convenient return paths.

Use gem trails, landmarks, sightlines, and lighting to guide exploration.

All 30 gems and all three dragons MUST be reachable.

Falls, broken barriers, and checkpoint transitions MUST NOT create softlocks.

VISUAL DIRECTION

Use a cohesive, colorful low-poly fantasy style:

- Grassy hills and pale stone architecture.
- Castle towers, archways, and colorful roofs.
- Warm sunlight, soft shadows, atmospheric fog, and a sky gradient.
- Distinct silhouettes for enemies, crystals, gems, and scenery.
- Rotating gems with sparkle and pickup effects.
- Fire particles aligned with the actual attack.
- Crystal fragments and magical rescue effects.
- A clear visual transformation when the portal unlocks.

Reuse geometry and materials where appropriate. Bound particle counts and clean up transient effects. Handle resizing correctly and keep rendering resolution reasonable on high-DPI displays.

DRAGON RESCUE SYSTEM

Place three adult dragons inside magical crystals.

Each dragon **MUST** have:

- An original name.
- A distinct color scheme.
- A recognizable personality.
- A short line of dialogue.

Rescue roles:

- Courtyard dragon: welcoming; explains the objective.
- Tower dragon: adventurous; comments on gliding.
- Garden dragon: observant; rewards exploration.

Either a valid fire attack or a horn charge breaks a crystal.

Required rescue sequence:

1. Break the crystal and reveal the dragon.
2. Mark the dragon rescued exactly once.
3. Update the rescue counter and activate a nearby safe checkpoint.
4. Freeze player movement, enemy simulation, and damage.
5. Continue presentation animations and show the dialogue panel.
6. Require a fresh Enter or Space press to dismiss dialogue.
7. Play a farewell gesture and magical disappearance.
8. Restore gameplay and camera control cleanly.

Only one rescue sequence may run at a time.

Repeated hits **MUST NOT** duplicate progress, dialogue, or checkpoint events.

Rescue animations must continue even while gameplay simulation is frozen.

COMBAT, HEALTH, AND RESPAWNING

- Player starts with three health points.
- Enemies patrol, detect the player, pursue within a bounded area, and attack nearby.
- Enemies cannot move through walls or walk off platforms.
- Both fire breath and horn charges can defeat every enemy.
- Damage removes one health point.
- Taking damage grants brief, visibly indicated invulnerability.
- Defeated enemies provide clear visual feedback.

Falling below the level or reaching zero health:

- Respawn at the latest checkpoint.
- Restore all three health points.
- Reset player velocity, attacks, and temporary movement states.

- Reposition the camera safely.
 - Preserve collected gems, rescued dragons, broken barriers, and defeated enemies.
- Restart and Play Again MUST reset the entire run.
- Persistence across browser reloads is not required.

GAME STATE AND PROGRESSION

Use explicit game modes:

- Title
- Playing
- Rescue dialogue
- Paused
- Victory

Centralize authoritative progress state.

Derive HUD counters and portal status from that state.

Use stable identifiers for gems, dragons, and enemies to prevent duplicate events.

The portal unlocks if and only if:

gemsCollected >= 20 AND dragonsRescued === 3

Before unlocking:

- Approaching the portal shows the remaining requirements.
- Entering it cannot trigger victory.

After unlocking:

- Its appearance changes visibly.
- Entering it triggers victory exactly once.
- The victory screen shows gems collected out of 30, dragons rescued out of three, and a Play Again button.

USER INTERFACE

- Start screen with title, objective, controls, and Play button.
- HUD with health, gem count, and dragon count.
- Contextual guidance near the first crystal and locked portal.
- Rescue dialogue showing the dragon's name and message.
- Escape pause menu with Resume and Restart.
- Readable layout at different window sizes.
- Pause on browser focus loss.
- Handle pointer-lock loss without trapping the user or leaving movement active.
- Resume from a pause during a rescue into the correct rescue state.

IMPLEMENTATION WORKFLOW

Follow the existing orchestrator's full workflow.

Phase 0: Intake and plan

- Inspect the workspace and existing progress artifacts.
- Resume unfinished work when appropriate.
- Present the plan, assumptions, roles, and expected artifacts.
- Obtain the plan approval required by our orchestrator instructions before delegation.

Phase 1: Architecture

- Use scout when unfamiliar code requires exploration.
- Delegate architecture to architect.
- Define contracts for input, player, camera, collision, world entities, progress, UI, and lifecycle.
- Define initialization, update, reset, and disposal responsibilities.
- Exactly one module owns the main animation loop.

Phase 2: Tickets

- Delegate ticket creation to tech-lead.
- Assign implementation waves based on actual dependencies.
- Establish the shared skeleton and contracts before parallel implementation.
- Concurrent tickets MUST NOT own the same files.

Phase 3: Implementation

- Delegate one ticket per coder run.
- First complete movement, camera, and collision on a test area.
- Next build the complete traversable level.
- Add gems, enemies, combat, rescues, checkpoints, and portal progression.
- Finish with UI, animation, and effects.
- Run independent coders concurrently when their prerequisites are ready.
- Review and integrate each wave before dependent work begins.

Phase 4: Quality

- Delegate integrated verification to tester.
- Return defects to coder with exact evidence and reproduction steps.
- Independently review results before accepting them.

Phase 5: Documentation

- Delegate documentation to documenter.
- Document the actual implementation and verified commands.

Phase 6: Deployment validation

- Delegate production build and local preview validation to deployer.
- Use Vite development or preview servers bound to localhost.
- Use browser tools actually available in the environment.
- Do not assume particular browser tool names or operating-system commands.
- Do not publish publicly.

Phase 7: Final report

- Report delivered behavior, completed tickets, actual checks, limitations, and artifact locations.

Use repository files for handoffs.

Maintain docs/STATUS.md and the ticket index after each phase.

Follow the existing orchestrator instructions for scoped commits.

Preserve unrelated user changes.

REQUIRED ARTIFACTS

Follow the repository's existing layout or the orchestrator's standard paths:

- docs/architecture/components.md
- docs/architecture/interfaces.md
- docs/architecture/deployment.md
- docs/adr/
- docs/tickets/README.md
- docs/tickets/T-*.md
- docs/STATUS.md
- reports/quality-report.md
- README.md
- docs/api.md describing internal module contracts; no HTTP API is required
- docs/runbook.md
- docs/design.md

- Deployment configuration or instructions
- reports/deployment-validation.md

VERIFICATION AND ACCEPTANCE

Run the production build, TypeScript checks, and meaningful automated tests.

Automated tests MUST verify:

- Collecting one gem repeatedly counts it once.
- Hitting a rescued crystal cannot duplicate rescue progress.
- Portal stays locked with 19 gems and three dragons.
- Portal stays locked with 20 gems and two dragons.
- Portal unlocks with 20 gems and three dragons.
- Respawn preserves progress and restores health.
- Restart resets the entire run.
- Pause and rescue dialogue prevent gameplay damage.
- Victory cannot trigger more than once.

When browser tools are available, verify the actual game:

- Start from the title screen.
- Walk, jump, glide, charge, and breathe fire.
- Check walls, platform edges, ceilings, and camera obstruction.
- Fall and respawn at a checkpoint.
- Defeat enemies with both attack types.
- Reach and rescue all three dragons.
- Collect at least 20 gems.
- Confirm portal requirements and complete the level.
- Test pause, resume, Restart, and Play Again.
- Inspect screenshots of normal gameplay and a rescue.
- Check browser errors and asset-loading failures.

A successful build is not proof of playable behavior.

A screenshot is not proof that progression works.

Report unavailable checks as “not verified,” with the reason.

Do not fabricate test results or performance measurements.

DEFINITION OF DONE

The player can start, explore, fight, collect gems, rescue all three dragons, unlock the portal, win, and restart.

Controls work consistently.

Required objectives are reachable.

Respawning preserves progress.

Restart clears progress.

There are no known progression blockers or placeholder gameplay systems.

Deliver the complete runnable game and a concise final report with:

- What was built.
- How to install and run it.
- Which OpenHands subagents were actually used.
- Completed and open tickets.
- Actual build, test, and browser verification results.
- Known limitations and unverified checks.
- Relevant artifact paths and commits.